

HEP Agent Simulation Framework

Code Documentation

May 5, 2026

Contents

1	Introduction	3
1.1	Overview	3
1.1.1	Scientific Goal	3
1.1.2	Agent-Based Modeling	3
1.2	Architecture	4
1.3	Installation	4
1.3.1	Requirements	4
1.3.2	Building the Project	5
1.4	Running the Simulation	5
1.4.1	Python UI (Recommended)	5
2	Developing New Modules	6
2.1	Overview	6
2.2	Path A: Independent Module Development	6
2.2.1	Step 1: Implement the Logic (Fortran)	7
2.2.2	Step 2: Register the Module ID	8
2.2.3	Step 3: Add the Dispatch Case	8
2.2.4	Step 4: Expose in the UI	8
2.2.5	Step 5: Rebuild & Test	8
2.3	Path B: Collaborative Workflow (Test Modules)	8
2.3.1	Step 1: Open the Test Module File	9
2.3.2	Step 2: Write Your Logic	9
2.3.3	Step 3: Rebuild & Test	10
2.3.4	Step 4: Iterate	10
2.3.5	Step 5: Hand Off to Daniel	10
2.4	Design Philosophy: Infrastructure vs. Science	10
2.5	Design Philosophy: Infrastructure vs. Science	10
2.5.1	1. Internal Infrastructure (Core Mechanics)	11
2.5.2	2. Interface Infrastructure (The Bridge)	11
2.5.3	3. Science (Simulation Modules)	11
2.6	Best Practices	12

3	Data Structures & Storage	13
3.1	Data Storage Logic	13
3.1.1	Design Choice: Hybrid Storage	13
3.1.2	The Agent Array	14
3.1.3	Spatial Indexing: The Grid	14
3.2	Module <code>mod_agent_world</code>	15
3.2.1	Agent Lookup: Hashmap	15
3.3	Compaction	16
3.4	Watershed Clustering	16
3.4.1	Algorithm	16
3.4.2	Configuration	16
4	Fortran Only Version	18
4.1	Building the Executable	18
4.1.1	Prerequisites	18
4.1.2	Build Command	18
4.2	Running the Simulation	19
4.2.1	Execution	19
4.2.2	Command Line Arguments	19
4.3	Configuration	19

Chapter 1

Introduction

1.1 Overview

The HEP Agent Simulation Framework (**ABM-Hescor**) is a high-performance simulation environment designed to model agent dynamics over large-scale geospatial grids (Human Evolution Potential - HEP data). The core simulation logic is written in modern **Fortran 2003/2008** for maximum performance, while the user interface and high-level control are implemented in **Python** (via PyQt5).

1.1.1 Scientific Goal

The present program is intended to simulate various aspects of the lives of prehistoric humans—such as Neanderthals and other human species. These aspects include movement, life cycle (reproduction and death), cultural development, and genetic inheritance. The goal is to compare the results of the simulation with archaeological data and then use the simulation data to fill gaps in the fossil record. In the past, a predecessor of this simulation already demonstrated, using climate data, how the settlement of Europe by modern humans might have taken place.

1.1.2 Agent-Based Modeling

The simulation is an agent-based model (ABM). This means that individual acting entities within the simulation are treated separately, and the agents' decisions are not aggregated but made individually for each agent. For example, rather than solving a partial differential equation (PDE) for population density, the model simulates millions of individual agents moving, reproducing, and interacting.

Note for Researchers: If you are primarily interested in the scientific modeling aspects (equations, biological rules, environmental interactions), the chapter "**Developing New Modules**" is the most important part of this document for you.

1.2 Architecture

The system follows a hybrid architecture (Fortran + Python):

- **Fortran Backend:** Handles all heavy lifting including:
 - Memory management for millions of agents (using `mod_agent_world`).
 - Spatial indexing (Grid) and Agent-ID mapping (Hashmap).
 - Simulation modules (Movement, Births, Deaths, Resources).
 - High-speed I/O for HEP NetCDF files.
- **Python Frontend:** Provides a user-friendly interface for:
 - configuring simulations.
 - Real-time visualization of agents and data on a 3D/2D globe.
 - Controlling the simulation loop (Start, Stop, Step).
 - Debugging and verification.

Interaction between the two layers is handled via `f2py` generated wrappers, exposing Fortran subroutines directly to Python as a compiled extension module (`mod_python_interface`).

1.3 Installation

The program requires Python, Fortran (`gfortran`), and `pip`.

1.3.1 Requirements

```
1 sudo apt install gfortran
2 sudo apt install python3 python3-venv python3-pip
```

Python Packages

Visualisation requires:

```
1 pip install pandas matplotlib pillow cartopy numpy imageio
```

1.3.2 Building the Project

To compile the Fortran extension and link it to Python:

1. Navigate to the project root.
2. Run the build script:

```
1      ./build_extension.sh
2
```

3. Verify that `mod_python_interface*.so` exists in the root directory.

1.4 Running the Simulation

1.4.1 Python UI (Recommended)

The preferred way to run the simulation, especially for development and debugging, is via the Python user interface:

```
1 python3 python/application.py
```

This opens a GUI that allows you to:

- Load config files and HEP maps.
- Start/Stop/Restart the simulation.
- Visualize agent positions in real-time.
- Select and verify active simulation modules.

Chapter 2

Developing New Modules

2.1 Overview

A “Module” in this framework is a self-contained behavior that can be applied to agents or to the grid (e.g., Movement, Aging, Mating). Modules are written in Fortran, registered in the interface layer, and toggled on/off from the Python UI.

There are **two ways** to develop new simulation logic, depending on your experience level and workflow preference:

Path A — Independent Development

You create a full module from scratch: write the Fortran subroutine, register it, wire it into the dispatch loop, and add it to the UI. Best for users comfortable with the full codebase.

Path B — Collaborative (Test Modules)

You write your logic directly in the pre-wired `test_modules.f95` file, test it, and send the file to Daniel. Daniel then turns your prototype into a proper, permanent module. Best for users who want to focus on the science without touching multiple files.

2.2 Path A: Independent Module Development

Use this path when you are confident editing multiple Fortran and Python files. Figure 2.1 shows the full workflow.




Figure 2.1: Path A: Creating a standalone module requires changes in four files.

2.2.1 Step 1: Implement the Logic (Fortran)

Create a new subroutine. There are two patterns depending on what your module operates on:

Agent-Centric Module

Runs once per living agent. The subroutine receives a pointer to the current agent.

```

1 subroutine my_agent_behavior(agent_ptr)
2   use mod_agent_world
3   implicit none
4   type(Agent), pointer, intent(inout) :: agent_ptr
5
6   ! Example: age-based death
7   if (agent_ptr%age > 3000) then
8     call agent_ptr%agent_dies(reason=6)
9   end if
10 end subroutine my_agent_behavior

```

Grid-Centric Module

Runs once per tick over the spatial grid. Receives the full world container.

```

1 subroutine my_grid_behavior(w)
2   use mod_agent_world
3   implicit none
4   class(world_container), target, intent(inout) :: w
5
6   integer :: gx, gy
7   do gy = 1, w%config%dlat_hep

```

```

8         do gx = 1, w%config%dlon_hep
9             ! operate on w%grid%cell(gx, gy)
10        end do
11    end do
12 end subroutine my_grid_behavior

```

2.2.2 Step 2: Register the Module ID

In `src/interfaces/python_interface.f95`, define a new integer constant. Pick the next available number:

```

1 integer, parameter :: MODULE_MY_BEHAVIOR = 17

```

2.2.3 Step 3: Add the Dispatch Case

In the `step_simulation` subroutine, add a case to the `select case` block:

```

1 case (MODULE_MY_BEHAVIOR)
2     ! For agent-centric:
3     call apply_module_to_agents(my_agent_behavior, t)
4     ! For grid-centric:
5     ! call my_grid_behavior(world)

```

2.2.4 Step 4: Expose in the UI

In `python/spawn_editor.py`, add your module to the `available_modules` dictionary (ID must match the Fortran constant):

```

1 self.available_modules = {
2     # ... existing modules ...
3     "My New Behavior": 17,
4 }

```

2.2.5 Step 5: Rebuild & Test

```

1 ./build.sh
2 python3 python/application.py

```

Add your module from the Spawn Editor and run the simulation to verify.

2.3 Path B: Collaborative Workflow (Test Modules)

Use this path when you want to **focus on writing your logic** without needing to touch the interface, build scripts, or registration code. Everything is already wired up for you.

Figure 2.2 shows the collaborative workflow.




Figure 2.2: Path B: Edit the test module, rebuild, test, and iterate until ready.

2.3.1 Step 1: Open the Test Module File

Open `src/simulation_modules/mod_test_modules.f95`. You will find two subroutines:

`test_module_agents` Runs your code on **every living agent**, once per tick.

`test_module_grid` Runs your code **once per tick** on the entire grid.

Choose the one that fits your use case (or use both).

2.3.2 Step 2: Write Your Logic

Each subroutine has a clearly marked section:

```

1 ! =====
2 ! DECLARE YOUR VARIABLES HERE
3 ! =====
4
5 ! =====
6 ! YOUR CODE HERE
7 ! =====

```

Important:

- Declare all variables you need at the top of the subroutine (after `implicit none`).

- You do **not** need to modify any config files.
- Use `print *`, `"..."` followed by `call flush(6)` for debug output.
- The file contains detailed comments listing all available agent fields, grid accessors, and example code.

2.3.3 Step 3: Rebuild & Test

```
1 ./build.sh
2 python3 python/application.py
```

In the Spawn Editor, add “**Test Module (Agents)**” or “**Test Module (Grid)**” to the active module list, then run the simulation. Check the terminal for your `print` output.

2.3.4 Step 4: Iterate

Repeat steps 2–3 until your results are correct. When you are satisfied, move on to step 5.

2.3.5 Step 5: Hand Off to Daniel

Send your modified `mod_test_modules.f95` file to Daniel. He will:

1. Review and optimize your code.
2. Extract it into a proper, standalone module with its own file and config parameters.
3. Merge it into the main codebase so it becomes a permanent, selectable module.

2.4 Design Philosophy: Infrastructure vs. Science

2.5 Design Philosophy: Infrastructure vs. Science

A core architectural principle of the HEP Agent Simulation Framework is the separation of **infrastructure** (background mechanics) from **science** (model logic and equations).

However, this boundary is fluid. In practice, functions generally fall into three distinct categories:

2.5.1 1. Internal Infrastructure (Core Mechanics)

These functions manage memory arrays, update spatial grid configurations, or handle raw processing. They belong directly to the core data structures (e.g., `mod_agent_world` or `mod_grid_id`) and contain **no scientific equations**. They exist solely to maintain the integrity of the computer state.

Examples of Internal Infrastructure:

- `setup_grid`: Allocates the memory arrays for the simulation.
- `place_agent_in_grid` / `remove_agent_from_cell`: Updates spatial indices when an agent moves.
- `get_male_from_cell`: A technical helper procedure to retrieve data from memory.

2.5.2 2. Interface Infrastructure (The Bridge)

These are foundational functions that bridge the gap between memory management and scientific rules. While they belong to the structural code (like `mod_agent_world`), they are typically **called by science functions** to enact a biological event formally within the system.

Examples of Interface Infrastructure:

- `generate_agent_born`: Validates parents, initializes genetic tracking, and spawns the structural **Agent** entity into the world memory array.
- `agent_dies`: Formally flags the agent as dead, increments mortality tracking counters, and performs structural cleanup.

2.5.3 3. Science (Simulation Modules)

Functions that determine *why*, *when*, or *how* an agent behaves based on biological, ecological, or physical rules. These belong in the **Simulation Modules** layer (the "science" layer). These routines contain the actual mathematical mortality/birth equations and then call the Interface Infrastructure to execute the structural changes.

Examples of Science Procedures:

- **reviewed_agent_motion**: Uses the Langevin equation, random diffusion, and spatial gradients to physically dictate where the agent moves next. It mathematically determines *how* the entity traverses the HEP surface before calling the lower-level **update_pos** infrastructure to register the move.
- **agent_above_water**: A scientific validation function that checks ecological boundaries (is the agent's coordinate currently designated as water?). If so, it returns true, which allows higher science functions to dictate a response (like death or bouncing back).

By structuring the code across these three tiers, researchers can freely modify the high-level scientific equations without risking memory leaks or index corruption.

2.6 Clustering and Spatial Analysis

The simulation supports several algorithms for identifying agent populations and density peaks.

2.6.1 K-Means Clustering

Core K-Means clustering is implemented in `mod_kmeans.f95`. **Note on Initialization**: To ensure uniform coverage and isolation of geographically remote populations, the framework utilizes a **Farthest-First Traversal (K-means++)** algorithm for centroid initialization. This prevents redundant clustering within a single high-density peak when multiple remote agent islands exist.

2.6.2 DBSCAN (Density-Based Spatial Clustering)

DBSCAN is available for discovering clusters of arbitrary shapes and identifying noise (outliers).

- **Epsilon (ϵ)**: The maximum distance between two samples for one to be considered as in the neighborhood of the other.
- **MinPts**: The number of samples in a neighborhood for a point to be considered as a core point.

2.6.3 Convex Hull Fill

To improve visual consistency, a **Convex Hull Post-Processing** step is applied to clustered grid cells. This fills internal holes and noise gaps within identified clusters, preventing cross-cluster contamination while maintaining structural integrity.

2.7 Visualization and Rendering

The Python visualization engine (`simulation.py`) handles the translation of Fortran grid data into interactive maps.

2.7.1 Rendering Architecture

- **Bifurcated Loops:** 2D and 3D visual rendering are structurally independent. The `on_sim_progress` method utilizes different mathematical paths for drawing flat grids (`ImageItem`) versus interpolating spherical terrains (`GLMeshItem`).
- **RGBA Matrices:** The backend simulation pipeline expects arrays formatted exactly as **4-channel RGBA** (R, G, B, Alpha). Boundary edges and UI markers are "burned" directly into these matrices inside Python to ensure rendering certainty.

2.8 Best Practices

- **Statelessness:** Keep modules stateless where possible. Store state in agent fields or in module-level `save` variables.
- **Grid Access:** Use `w%grid` to access spatial neighbour information.
- **Performance:** Avoid file I/O inside per-agent loops. Use `print` sparingly (e.g., only on specific ticks).
- **Debug Output:** Always call `flush(6)` after `print` statements to ensure output appears before a potential crash.

Chapter 3

Data Structures & Storage

3.1 Data Storage Logic

Efficiently storing and retrieving agents is critical for performance. The framework uses a **Structure-of-Arrays (SoA)** approach mixed with object-oriented containers to balance cache efficiency and code maintainability.

3.1.1 Design Choice: Hybrid Storage

The simulation needs to handle millions of agents. Storing each agent as a separate object scattered in memory (as in a pure linked list) would cause excessive cache misses. Conversely, a pure structure-of-arrays can make code harder to read. We compromise by using a `world_container` that holds global arrays, but we expose individual agents via pointers to derived `Agent` types that "view" the array data.




Figure 3.1: Visual representation of the hybrid data structure. The `world_container` acts as the root, managing contiguous memory arrays for agents while maintaining spatial (Grid) and logical (Hashmap) indices for efficient access.

3.1.2 The Agent Array

Agents are stored in a large, pre-allocated 2D array within the `world_container`:

```
1 type(Agent), allocatable :: agents(:, :) ! (max_size, npops)
```

- **Rows:** Iterate through individual agents (up to `max_size`).
- **Columns:** Separate different populations (e.g., `npops`).

This allows for contiguous memory access when iterating over a single population, which is the most common operation.

3.1.3 Spatial Indexing: The Grid

For spatial interactions (e.g., mating, conflict), agents are indexed by their geospatial location using the **Grid** structure.

- Each grid cell (`grid_cell`) contains a dynamic list of Agent IDs currently located in that cell.
- When an agent moves, it is removed from the old cell's list and added to the new cell's list.

Module Types: Agent-Centric vs Cell-Centric

When developing modules, you must decide the primary access pattern:

- **Agent-Centric Modules:** Functions that apply to a specific agent (e.g., Movement). These iterate over the `agents` array.

```
1     subroutine agent_move(agent_ptr)
2         type(Agent), pointer :: agent_ptr
3     end subroutine
```

- **Cell-Centric Modules:** Functions that apply to a location (e.g., Overpopulation Death). These iterate over the `grid`.

```
1     subroutine death_overpopulation(cell_ptr)
2         type(cell), pointer :: cell_ptr
3     end subroutine
```

3.2 Global State & Metrics

In addition to individual agent data, the `world_container` tracks macroscopic simulation states and per-tick accumulators to monitor population health and system performance.

3.2.1 Tick Accumulators (`t_tick_accumulators`)

A history of accumulators is maintained (`accumulators_history(:)`) to track events within a single tick. These are automatically reset at the start of each step:

- `phi_birth_acc` / `phi_death_acc`: Cumulative birth and death counts.
- `n_alive_acc`: Count of agents currently alive.
- `ms_per_tick`: Execution time for the current simulation step.

3.2.2 Dynamic State (`t_dynamic_state`)

Global variables that evolve over time and influence agent behavior:

- `K_fertility`: A macroscopic scaling factor used to regulate population growth according to the Verhulst model.

3.3 Module `mod_agent_world`

This module defines the central `world_container` which holds:

- `agents`: The main agent array.
- `grid`: The spatial grid.
- `index_map`: The ID lookup hashmap.
- `config`: Global configuration parameters.
- `accumulators_history`: Buffer of recent performance and biological metrics.
- `dynamic_state_vars`: Current macroscopic simulation state.

It serves as the "Hub" for all data access.

3.3.1 Agent Lookup: Hashmap

To find an agent by its unique ID (without searching the entire array), the system uses a custom integer hashmap (`mod_hashmap`).

- **Key**: Agent ID (Integer).
- **Value**: A composed integer packing both `population_index` and `array_index`.

This allows $O(1)$ access to any agent pointer given its ID.

3.4 Compaction

When agents die, they are marked with `is_dead = .true..` They are NOT removed immediately to avoid array resizing costs. Instead, a `compact_agents` subroutine is called periodically (e.g., end of timestep).

1. It iterates from both ends of the array.
2. Dead agents on the left are overwritten by living agents from the right.
3. The array size is logically reduced (tracking integer), but physically remains allocated.

3.5 Density Analysis & Clustering

The simulation includes a suite of spatial analysis modules (`mod_watershed`, `mod_kmeans`, `mod_clustering`) that group grid cells into clusters based on agent density and HEP suitability.

3.5.1 Global Density Smoothing

To ensure consistency across different clustering algorithms, a unified density surface (`human_density_smoothed`) is generated once per tick. This surface is processed using an iterative box filter.

3.5.2 Algorithms

1. **Watershed**: Performs gradient ascent on the global density surface.
2. **K-Means++**: Utilizes a **Farthest-First Traversal** initialization to ensure centroids are placed in remote agent populations, avoiding redundant clustering on a single peak.
3. **Convex Hull Fill**: A post-processing step that fills internal holes in clusters, ensuring structural integrity for visualization and demographic analysis.

3.5.3 Configuration

Clustering parameters are defined in the `namelist /config/` block:

Parameter	Type	Default	Description
<code>human_density_smoothing_radius</code>	integer	2	Radius of the box-filter applied to the global density surface.
<code>human_density_smoothing_iterations</code>	integer	1	Number of iterative passes for the box-filter (high-density smoother).
<code>watershed_threshold</code>	real(8)	0.05	Minimum density for a cell to be included in a cluster.

Chapter 4

Fortran Only Version

The simulation framework includes a standalone Fortran executable, allowing users to run the model without the Python interface or GUI. This is particularly useful for high-performance computing (HPC) environments, batch processing, or profiling.

4.1 Building the Executable

The standalone version is built using the provided shell script `build_fortran.sh`. This script compiles all necessary modules, creates a static library (`libhep.a`), and links the final executable.

4.1.1 Prerequisites

- **Compiler:** `gfortran` (or compatible Modern Fortran compiler).
- **Libraries:** NetCDF and NetCDF-Fortran development libraries (`libnetcdf-dev`, `libnetcdf-f-dev`).

4.1.2 Build Command

To build the project, run the following command from the project root:

```
1 ./build_fortran.sh
```

This will create:

1. A `build/` directory containing compiled object files (`.o`) and module files (`.mod`).
2. A static library `build/libhep.a`.
3. The standalone executable `build/main_fortran`.

4.2 Running the Simulation

The executable can be run directly from the command line. It reads the specific configuration file located at `input/config/main_fortran_config.nml`.

4.2.1 Execution

```
1 ./build/main_fortran
```

4.2.2 Command Line Arguments

The standalone version supports command-line arguments to override specific runtime parameters:

- `-output_interval <integer>`: Sets the interval (in ticks) at which simulation status is printed to the console. Default is 1000.

Example:

```
1 ./build/main_fortran --output_interval 500
```

4.3 Configuration

Unlike the Python version which uses `basic_config.nml`, the standalone Fortran version hardcodes the configuration path to:

`input/config/main_fortran_config.nml`

Ensure this file exists and contains valid namelists (`&dims` and `&config`) before running. The structure of this file is identical to the standard configuration files used by the GUI.